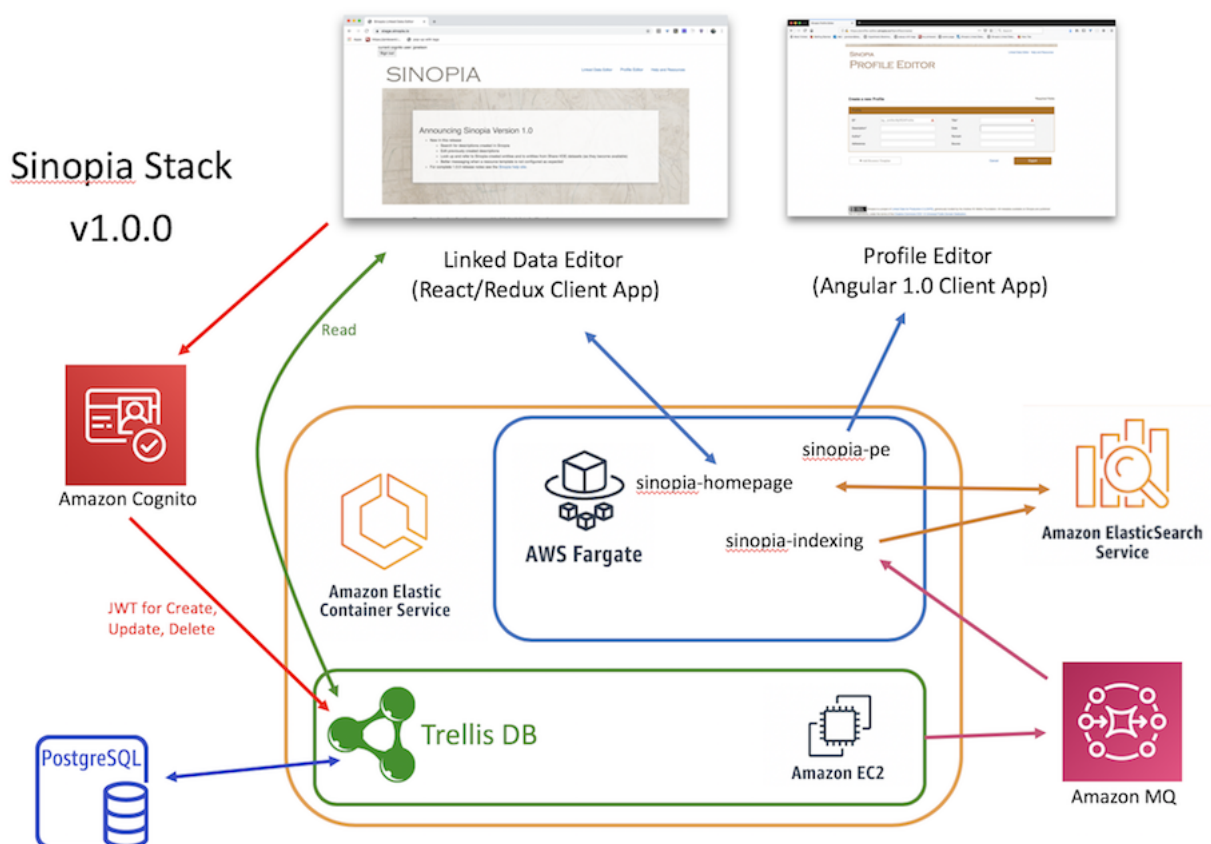




2019 BIBFRAME Workshop - Stockholm, Sweden

<https://ld4p.github.io/bf-workshop-2019/sinopia-stack-aws/>

Running the Sinopia Stack on Amazon Web Services



Background

An early requirement for the open-source [Sinopia](#) project was to build a cloud-based collaborative editing environment. The team chose [Amazon Web Services](#), one of the most popular commercial cloud provider, by utilizing a number of different AWS services and products to host Sinopia and its dependent technologies. Although the first version of Sinopia is closely tied to specific AWS cloud services, we are looking at more generic infrastructure options for hosting Sinopia on other commercial cloud services as well as in hybrid environments.

Treillis - Linked Data Platform

In Sinopia's early analysis, we determined that having a RDF triplestore was not necessary for meeting the requirements of a create-update-read-delete (CRUD) editor for RDF.



Amazon Web Services

Cognito

To handle authentication and authorization for create, update, and delete operations on the Linked Data RDF and JSON resource templates in Sinopia, we are using the AWS [Cognito](#) secure user sign-up and access control service.

Amazon Web Services

Elasticsearch Service

After a user successfully signs-up via the [Amplify](#) SDK, a [JSON Web Token](#).

[Top](#)

For the initial 1.0.0 release, a simple search index is indexed and searched through the AWS hosted [Elasticsearch Service](#). Elasticsearch is a full-text search engine, based on [Lucene](#), that has been optimized for running on Amazon's cloud.



Amazon Web Services

Elastic Container Service

The deployment of Sinopia on AWS relies on pre-built Docker images hosted on [DockerHub](#) that are then run in a [Elastic Container Service](#) (ECS) cluster. Sinopia is run on three [ECS](#) clusters; development, staging, and production with corresponding Docker images for each environment.



Amazon Elastic Container Service

[Top](#)



Amazon Web Services

Apache ActiveMQ

Trellis publishes events like creating, updating, or deleting resource as they occur to a AWS [Messaging Queue](#) that is monitored by a process called `sinopia-indexing` pipeline that then updates the Elasticsearch search index.

```
message = {
  payload: '2345'
}
```

[Top](#)



Amazon Web Services

Relational Database Service (RDS)

Sinopia server's foundation is a variant of [Trellis](#) Linked Data Platform that uses a [PostgreSQL](#) relational database to manage the Sinopia Editor's RDF and JSON payloads instead of a RDF Triplestore.

The AWS RDS services offers the following advantages:

- Significantly cheaper than AWS Triplestore [Neptune](#)
- Automatic backups
- Administrative overhead reduced

[Top](#)

Amazon Web Services

Fargate



[Fargate](#) is an AWS service that runs [Docker](#) containers in a state-less fashion with in Virtual Elastic cluster. A [Docker](#) is a lightweight running linux process that is generated in a deterministic way from a [Docker](#) image


In Sinopia, we run the following Docker Images as [Fargate](#) tasks:

- Sinopia Linked Data Editor [Docker image](#)
- Sinopia Profile Editor [Docker image](#)
- Sinopia Indexing Pipeline [Docker image](#)
- Sinopia ACL [Docker image](#)

[Top](#)

Amazon EC2

The other method we use to run a [Docker](#) image is by hosting as part of a virtual machine running as a [EC2](#) Docker container.



When we started Sinopia, [Fargate](#) does not persist data so we could not mount a permanent disk volume to store the RDF entities [Memento](#) metadata as needed by our Linked Data Platform, [Trellis](#).

[Top](#)

©2019 Jeremy Nelson under the [CC4](#) license.